

# Logistics Delivery Time Prediction (ridge regression l2 norm)

---

## Training Summary

Model Type	Feedforward Neural Network (Keras Sequential)
Epochs	15
Optimizer	Adam
Loss Function	Mean Squared Error
Regularizer	L2 (Ridge) with factor 0.01
Final Test MAE	7.04 minutes

## Prediction Examples

Actual Time (min)	Predicted Time (min)	Absolute Error (min)
20.0	22.43	2.43
40.0	35.77	4.23
51.0	50.99	0.01
41.0	41.47	0.47
61.0	54.23	6.77

## Notes on Model and Regularization

- The model is a feedforward neural network trained to predict delivery times.
- L2 regularization (also known as Ridge regression in linear models) is applied to the network weights via the `kernel_regularizer` parameter.
- L2 regularization helps to reduce overfitting by penalizing large weights, improving model generalization.
- This approach is similar in principle to Ridge regression, but applied here to a deep learning model.

## Mean Absolute Error (MAE)

MAE measures the average magnitude of errors in predictions, without considering their direction.

It is calculated as:

$$MAE = (1/n) * \sum |y_i - \hat{y}_i|$$

where:

-  $y_i$  is the actual value,

- $\hat{y}_i$  is the predicted value,
- $n$  is the number of samples.

## Conclusion

- The model achieved a test Mean Absolute Error (MAE) of approximately 7.04 minutes, indicating good predictive performance.
- Predictions are close to actual delivery times, suitable for operational use.
- Further feature engineering or additional data may improve accuracy.