

Ridge(l2)_Lasso(l1)_ElasticNet(l1+l2) Regression model

Standardized Evaluation Metrics Table

Model	Alpha	l1_ratio	MSE	RMSE	MAE	R ²
Ridge	0.0001	–	8337.40	91.31	46.01	0.9713
Lasso	0.5964	–	8367.80	91.48	45.35	0.9712
Ridge (CV Tuned)	0.0001	–	8013.62	89.52	45.31	0.9724
Lasso (CV Tuned)	0.5964	–	8360.85	91.44	45.48	0.9712
ElasticNet (CV Tuned)	0.0054	0.9	8350.80	91.38	46.07	0.9712

Model Comparison Summary

- **Best Overall Model:**
Ridge (CV Tuned) achieved the lowest MSE (8013.62) and the highest R² (0.9724), indicating the best fit and generalization performance on your test data.
- **Lasso vs ElasticNet:**
Both models perform similarly, with Lasso having a slightly lower MAE and ElasticNet a slightly lower MSE, but neither surpasses Ridge.
- **Base Ridge vs Tuned Ridge:**
Ridge performance greatly improves after tuning, demonstrating the importance of cross-validation.

Hyperparameter Summary

- **Ridge α = 0.0001** → Small regularization strength, indicating a low-bias model performs well.
- **Lasso α = 0.5964** → Stronger regularization; selects fewer features.
- **ElasticNet α = 0.0054, $l1_ratio$ = 0.9** → Mostly Lasso-like behavior.